




BOURBAKI

Ernesto Lupercio  
May 6, 2025

# Understanding GRPO



## What Is REINFORCE?

- REINFORCE is a way to teach a computer (or robot) to learn from rewards.
  - It tries things, sees what works, and adjusts itself to do better over time.
  - Like learning to play a game by trial and error.
- 

## How Does It Work?

- The computer sees a state  $s$  (like the game screen).
- It chooses an action  $a$  (like jump or move).
- It gets a reward  $R$  (like a score).
- It uses this info to adjust its brain  $\theta$ .

## What Is a Policy?

- A policy tells the computer what to do in each situation.
- Written as  $\pi_\theta(a | s)$ : the chance of taking action  $a$  in state  $s$ .
- $\theta$ : the computer's internal settings (like its brain weights).
- Changing  $\theta$  changes how it plays.

## Learning Rule (Simplified)

- After each move, it adjusts its brain based on how good the result was:

$$\theta \leftarrow \theta + \text{learning rate} \times (\text{action score})$$

- Good actions  $\Rightarrow$  increase chance
- Bad actions  $\Rightarrow$  decrease chance

## Mathematical Form (Optional)

$$\theta \leftarrow \theta + \alpha \cdot \nabla \theta \log \pi_{\theta}(a_t | s_t) \cdot R_t$$

- $\alpha$ : learning rate
- $\nabla \theta$ : change in settings
- $\log \pi_{\theta}(a_t | s_t)$ : how confident it was
- $R_t$ : the reward it got

## Why It's Called REINFORCE

- Like a coach saying "Great move! Do that again."
- Or "That was bad. Let's change the plan."
- It reinforces good actions with higher probability.

## Everyday Analogy

- Imagine throwing darts at a target.
- You get points based on where you hit.
- You adjust your next throw based on the last result.
- Over time, you learn to aim better.





# Summary

- REINFORCE teaches by trying and adjusting.
- It uses feedback (rewards) to get better.
- It increases good decisions and decreases bad ones.
- It's the starting point of many powerful AI systems today.



## REINFORCE : Key Formula

$$\nabla_{\theta} J(\theta) = E[\nabla_{\theta} \log \pi_{\theta}(a_t|s_t) R_t]$$

- $\theta$ : parameters of the policy (like the robot's brain settings)
- $\pi_{\theta}(a_t|s_t)$ : probability of taking action  $a_t$  in state  $s_t$
- $R_t$  : total reward after time  $t$
- $\nabla_{\theta} \log \pi_{\theta}(a_t|s_t)$ : how to change the policy

## What Does It Mean?

- Increase the chance of actions that gave good rewards.
- Decrease the chance of actions that led to bad outcomes.
- Learn by trial and error using total reward.

## Symbols in Context

- State ( $s_t$ ): What the agent sees at time  $t$
- Action ( $a_t$ ): What the agent decides to do
- Reward ( $R_t$ ): What it gains (or loses) afterward
- Policy ( $\pi_\theta$ ): Rule the agent follows to choose actions
- Gradient ( $\nabla \theta$ ): Tells how to adjust the brain

## Pseudocode for REINFORCE

Initialize policy parameters  $\theta$

Repeat:

Generate an episode:  $(s_0, a_0, r_1, \dots, s_T)$

For each time step  $t$  in episode:

Compute return  $R_t = \text{sum of future rewards}$

Compute gradient:  $g = \text{grad}_{\theta} \log(\pi_{\theta}(a_t | s_t)) * R_t$

Update  $\theta$  using gradient ascent

Until convergence

# Actor-Critic

- Teach a robot (or computer) to make good decisions.
- Use rewards to learn what works well.
- Actor-Critic is a smarter way to learn than REINFORCE.

## Two Roles: Actor and Critic

- Actor: decides what to do (takes actions).
- Critic: gives feedback (estimates how good the situation is).
- They work together to learn faster and better.

## Key Idea

- Actor learns the policy  $\pi_{\theta}(a|s)$ : what to do in state  $s$ .
- Critic learns the value  $V_w(s)$ : how good is state  $s$ ?
- Together, they adjust based on feedback.



## Mathematical Pieces

- Actor uses the policy gradient:

$$\nabla_{\theta} J(\theta) = \nabla_{\theta} \log \pi_{\theta}(a|s) \cdot A(s, a)$$

- Critic estimates the advantage:

$$A(s, a) = r + \gamma V(s') - V(s)$$

- $\gamma$ : discount factor (between 0 and 1)

## How the Critic Helps

- Instead of waiting until the end of the game (like REINFORCE),
- The critic gives feedback immediately using its value estimate.
- This helps the actor make better choices faster.

## Learning Together

1. Actor takes action  $a$  in state  $s$ .
2. Environment gives reward  $r$  and next state  $s'$ .
3. Critic updates  $V(s)$  using:  
$$V(s) \leftarrow V(s) + \alpha(r + \gamma V(s') - V(s))$$
4. Actor updates policy based on advantage.

## Why Is This Better Than REINFORCE?

- Less randomness (variance)
- Learns from every step, not just full games
- Actor doesn't need to guess — critic helps guide it

## Everyday Analogy

- Actor = student trying to answer questions
- Critic = teacher giving hints and feedback
- Together, they help the student (robot) learn faster

# Conclusion

- Actor-Critic uses two brains: one to act, one to evaluate.
- It builds on REINFORCE but is more efficient.
- A powerful idea used in many modern AI systems.



## TRPO (Trust Region Policy Optimization)

- Imagine training a robot that suddenly forgets everything after one mistake.
- Regular policy gradient can make big, unsafe changes.
- We need a way to say: "Don't change the brain too much in one go."
- That's what TRPO does: it keeps updates safe and small.



## The TRPO Idea

- TRPO updates the policy, but only if it doesn't stray too far.
- Think of it like driving: you can steer, but not jerk the wheel.
- Mathematically, it uses a constraint:
$$\text{KL}(\pi_{\text{old}} // \pi_{\text{new}}) \leq \delta$$
- KL divergence measures how different the two policies are.



## The TRPO Formula

- TRPO maximizes a surrogate reward function:

$$L(\theta) = E \left[ \frac{\pi_{\theta}(\mathbf{a}|\mathbf{s})}{\pi_{\theta_{\text{old}}}(\mathbf{a}|\mathbf{s})} A(\mathbf{s}, \mathbf{a}) \right]$$

- This compares the new and old policies.
- But only allows it if the KL difference is small.

## Why This Matters

- It prevents the robot from making dangerous changes.
- Every update is like a gentle nudge, not a hard shove.
- Result: safer learning, more stable improvement.



## The PPO Idea

- Instead of hard rules (like  $KL \leq \delta$ ), PPO softens it.
- It compares new vs. old action probabilities:

$$r_t(\theta) = \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\theta_{\text{old}}}(a_t | s_t)}$$

- If the change is too big, we clip it to keep it safe.

## Why PPO Was Invented

- TRPO works well, but is complicated to implement.
- PPO keeps the same goal: avoid dangerous updates.
- But it uses a much simpler trick: clipping.
- PPO is like saying: “Let the robot learn, but not too fast.”

## The PPO Objective

- PPO tries to maximize:
$$\min (rt \text{ } At, \text{clip}(rt, 1 - \epsilon, 1 + \epsilon)At )$$
- This means:
  - If the update is small, let it happen.
  - If it's too big, clip it so it doesn't go wild.
- This makes PPO robust and stable.

## Understanding the clip() Function

- PPO uses a clipping function to avoid large policy updates.
- Mathematically:

$$\text{clip}(r, 1 - \epsilon, 1 + \epsilon) = \begin{cases} 1 - \epsilon & \text{if } r < 1 - \epsilon \\ r & \text{if } 1 - \epsilon \leq r \leq 1 + \epsilon \\ 1 + \epsilon & \text{if } r > 1 + \epsilon \end{cases}$$

- In words: "Don't let the ratio  $r$  go too far from 1."
- Keeps the policy update from being too aggressive.
- Helps maintain stability and trust in learning.

## PP0 Symbols Explained

- $\pi_\theta$  : policy with parameters  $\theta$
- $a_t$  : action taken at time  $t$
- $s_t$  : state observed at time  $t$
- $A_t$  : advantage at time  $t$ , how much better the action was than average

$r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)}$ : how much the new policy favors this action

$\epsilon$ : clip range parameter, e.g., 0.2

## Pseudocode for PPO

```
Initialize policy parameters  $\theta$  and value function parameters  
for iteration in range(N):  
    Collect trajectories using current policy  $\pi_{\theta}$   
    Estimate advantage  $A_t$  and returns  
    for epoch in range(K):  
        for minibatch in trajectories:  
             $r_t = \pi_{\theta}(a_t|s_t) / \pi_{\theta_{old}}(a_t|s_t)$   
             $clipped = clip(r_t, 1 - \epsilon, 1 + \epsilon)$   
             $loss = -\min(r_t * A_t, clipped * A_t)$   
        Update policy parameters using gradient of loss  
    Update value function to fit returns
```



## Why PPO Works Well

- Easier to implement than TRPO.
- Works with simple gradient descent (no fancy math).
- Learns safely and quickly.
- PPO is used in real-world tasks — from games to finance.

## History and Impact of PPO

- Proposed by OpenAI in 2017. Aimed to simplify and stabilize policy optimization.
- Built as a lighter alternative to TRPO (Trust Region Policy Optimization).
- Became widely used for its balance of performance and ease of implementation.
- Published in: "Proximal Policy Optimization Algorithms", Schulman et al., 2017.

## Spectacular Applications of PPO

- **OpenAI Five** — trained agents to play Dota 2 at a superhuman level.
- **Robotic control** — simulated and real robots for walking, grasping, balancing.
- **Autonomous vehicles** — learning lane-keeping and adaptive driving policies.
- **Finance** — portfolio management, algorithmic trading strategies.
- **Games** — used in Unity ML-Agents Toolkit to train AI in video games.

## Group Relative Policy Optimization (GRPO): Origins and Development

- **Introduction of GRPO:** Introduced by DeepSeek in their 2024 paper, DeepSeekMath: Pushing the Limits of Mathematical Reasoning in Open Language Models [?].
- **Motivation:** Designed to enhance mathematical reasoning in large language models (LLMs) by simplifying the reinforcement learning (RL) process.
- **Key Innovation:** Eliminates the need for a separate value function (critic) by estimating baselines from group scores, reducing computational overhead.
- **Implementation:** First applied in training DeepSeek-R1 and DeepSeek-R1-Zero models, demonstrating improved reasoning capabilities without extensive supervised fine-tuning.

## DeepSeek-R1-Zero: Advancing LLM Reasoning with GRPO

- **Model Overview:** DeepSeek-R1-Zero is a large language model trained exclusively using GRPO, without supervised fine-tuning, achieving strong performance in reasoning tasks.
- **Training Efficiency:** Leveraged GRPO to reduce the need for human-annotated data and extensive computational resources, making training more cost-effective.
- **Performance:** Demonstrated competitive results on benchmarks like the American Invitational Mathematics Examination (AIME) and MATH, showcasing GRPO's effectiveness in enhancing LLM reasoning abilities.
- **Impact:** DeepSeek-R1-Zero's success with GRPO has influenced the development of subsequent models and training methodologies in the field of AI.

## GRP0: Learning Without a Teacher

Imagine you're teaching a robot to solve math problems.

- You give the robot a math question.
- It tries out  $k$  different answers:  $a_1, a_2, \dots, a_k$ .
- But there's no teacher around to say which is correct!
- So the robot compares the answers against each other.
- The ones that look better get rewarded, the rest don't.

## How Does GRPO Know Which Answer Is Better?

- Even without a teacher, GRPO can compare answers by scoring them using a reward model.
- Think of it as an internal judge trained to recognize good answers.
- For example, it might score answers based on:
  - Correctness clues (e.g., algebraic validity)
  - Completeness and clarity
  - Alignment with the question asked
- Then it ranks all answers  $a_1, \dots, a_k$  by these scores.
- The best answers (relative to the rest) get positive rewards  $R_i > R^-$ .
- GRPO uses these scores to update the robot's brain — favoring better answers.

## GRPO: The Update Rule

The robot learns by making better answers more likely.

- Each answer gets a relative score — how good it was compared to the rest.
- Define:

$$\theta \leftarrow \theta + \alpha \sum_{i=1}^k \nabla_{\theta} \log \pi_{\theta}(a_i | s) \cdot (R_i - \bar{R})$$

Where:

- $\theta$ : the robot's brain (model weights)
- $\alpha$ : learning rate (how fast it learns)
- $\pi_{\theta}(a_i | s)$ : the probability it picked answer  $a_i$
- $R_i$ : reward for answer  $a_i$
- $\bar{R}$ : average reward of all answers



## GRP0: All Symbols Explained

- $s$ : the question or input the robot sees
- $a_i$ : the  $i$ -th answer the robot gives
- $\pi_{\theta}(a_i | s)$ : the probability assigned to that answer
- $R_i$ : a reward for how good that answer was (relative to others)

$$\bar{R} = \frac{1}{k} \sum_{j=1}^k R_j: \text{ the average reward over all answers}$$

- $\nabla_{\theta} \log \pi_{\theta}(a_i | s)$ : tells how to increase or decrease the chance of this answer
- $\alpha$ : how big of a change to make (learning rate)

## GRPO Pseudocode

for each question  $s$ :

Generate  $k$  answers  $a_1, \dots, a_k$  from current policy  $\pi_\theta$

Score each answer using a reward model:  $R_1, \dots, R_k$

Compute average reward:  $R_{\text{bar}} = \text{average}(R_i)$

For each  $i = 1$  to  $k$ :

Compute  $\log_{\text{prob}} = \log \pi_\theta(a_i | s)$

Compute gradient  $= (R_i - R_{\text{bar}}) * \text{grad}_\theta \log_{\text{prob}}$

Accumulate gradient

Update  $\theta$  using accumulated gradient

## GRPO: Why It Works


- If an answer is better than average:  $R_i > \bar{R} \Rightarrow$  increase chance of that answer.
- If an answer is worse: decrease its chance.
- No need for labeled answers — the robot self-improves.
- It's like playing chess with yourself and always favoring better moves.


## DeepSeek and GRPO in Practice

- DeepSeek used GRPO to train a math-solving robot called
- DeepSeek-R1-Zero.
- It didn't get supervised feedback — just learned from its own output quality.
- GRPO made training faster and cheaper — no human labels needed.
- It ended up solving hard math problems like those in international competitions!

# Why GRPO is Awesome

- Self-taught learning: the robot becomes smarter by comparing its own attempts.
- No teachers needed: works even when no labeled data is available.
- Efficient and powerful: used in cutting-edge AI research.
- Good for logic, math, puzzles, and reasoning tasks.





BOURBAKI

COLEGIO DE MATEMÁTICAS